
CS335 Class Notes

Luis F. Uceta

2020-01-28

Contents

1	Lecture 1: C++ basics	4
1.1	C++ classes	4
1.2	Initialization list	4
1.3	Constant member functions	5
1.4	Interface and implementation	5
1.4.1	IntCell interface	6
1.4.2	IntCell implementation	6
1.4.3	Using IntCell	7
1.5	C++ Objects	7
1.6	vectors and strings	7
1.6.1	vector	8
1.7	C++ details	8
1.7.1	Pointers	8
1.7.2	Dynamic object allocation and garbage collection	9
1.7.3	Lvalues, and Rvalues	9
1.7.4	References	10
1.7.5	Uses of lvalue references	10
2	Lecture 2: C++ Basics (cont.)	12
2.1	Parameter passing	12
2.2	Return passing	13
2.3	C++11 <code>std::swap</code> and <code>std::move</code>	14
2.4	The big five	15
2.4.1	Destructor	15
2.4.2	Copy constructor and move constructor	16
2.4.3	Copy assignment and move assignment (operator =)	16
2.4.4	Defaults	17
2.5	Templates	18
2.5.1	Function templates	18
2.6	Function objects	19
2.6.1	Class templates	23
2.7	Using matrices	23
2.8	Class problems	25
3	Lecture 3: Algorithm analysis	27
3.1	Mathematical background	27

3.2	Notation	27
3.2.1	Big-Oh notation	27
3.2.2	Big-Omega notation	28
3.2.3	Big-Theta notation	29
3.3	Typical growth rates	29
3.4	Rules	30
3.5	Few points	30
3.6	Model of computation	31
3.7	What to analyze	31
3.7.1	Types of performance	31
3.8	Running-time calculations	32
3.9	Sample problems	33
3.9.1	Sum of cubes	33
3.9.2	Factorial	33
3.9.3	Maximum subsequence sum problem	34
3.9.4	Algorithm 1 (brute force)	35
3.9.5	Algorithm 2 (brute force)	36
3.9.6	Algorithm 3	36
3.9.7	Algorithm 4	37
4	Lecture 4: Algorithm analysis (cont.) and lists/stacks/queues	38
4.1	Binary search	38
4.1.1	Running time	38
4.2	List/stacks/queue ADTs	39
4.2.1	The List ADT	39
4.2.2	Vector in the STL	40
4.2.3	List in the STL	40
4.3	Containers	41
4.4	Iterators	41
4.4.1	How to get an iterator	41
4.4.2	Iterator methods	41
4.4.3	Container operations that require iterators	42
4.4.4	Example: Using erase on a list	43
4.4.5	const_iterators	43
4.4.6	Printing a container	44

1 Lecture 1: C++ basics

1.1 C++ classes

```
class IntCell {
public:
    // This is a constructor
    IntCell() { _value = 0; }

    // This is another constructor. It takes an initial value.
    IntCell( int initialValue ) { _value = initialValue; }

    // This is an accessor (or getter) method.
    int read() { return _value; }

    // This is a setter method.
    void write( int newValue ) { _value = newValue; }

private:
    int _value;
};

int main() {
    // Using the first constructor w/o parameters
    IntCell x;
    int xValue = x.read(); // 0
    x.write(5);
    xValue = x.read();    // 5

    // Using the first constructor with a single parameter
    IntCell y(12);
    int yValue = y.read(); // 12
    y.write(y.read() * 2);
    yValue = y.read();    // 24
}
```

The keywords **public** and **private** determine the visibility of class members. A member that is **public** may be accessed by any method in any class while a member that is **private** may only be accessed by methods in the class it's declared. Data members are usually declared **private** to hide a class's internal details (i.e., **information hiding**).

1.2 Initialization list

Instead of initializing data members inside a constructor's body, an **initialization list** can be used to do so right in the constructor's signature:

```
class IntCell {  
    public:  
        explicit IntCell( int initialValue = 0 ) : _value { initialValue }  
        { }  
        int read() const { return _value; }  
        void write( int newValue ) { _value = newValue; }  
  
    private:  
        int _value;  
};
```

By default C++ does behind-the-scenes type conversions in all one-parameter constructors. Thus, to avoid this the keyword `explicit` is used.

With **implicit type conversion**,

```
IntCell obj;  
obj = 37;
```

means

```
IntCell temp = 37;  
obj = temp;
```

With **explicit type conversion** (by using `explicit`), a “type mismatch” compiler error is thrown instead. The use of `explicit` means a one-parameter constructor cannot be used to generate an implicit temporary object.

1.3 Constant member functions

A member function that examines but does not change the state of its object is an **accessor**. In the `IntCell`, `read` is an accessor and thus it's marked explicitly as an accessor by using the keyword `const` after the closing parenthesis that ends the parameter list. This signals that this method doesn't change the state of an `IntCell` and if it tries the compiler complains.

1.4 Interface and implementation

An **interface** lists the class and its members (both data and methods) and the **implementation** provides the implementations of the methods.

In C++, this is done by placing the interface in a `.h` file and the implementation in a `.cpp` file. Source code that requires knowledge of the interface must include the interface file (e.g., `#include Interf`

.h). To avoid including files multiple times, a few preprocessor commands are used in the interface file:

```
#ifndef INTCELL_H // if INTCELL_H is not defined...
#define INTCELL_H // ...define it.

// Here's the IntCell class's interface

#endif // end the definition
```

1.4.1 IntCell interface

```
// IntCell.h
#ifndef INTCELL_H
#define INTCELL_H

class IntCell {
public:
    explicit IntCell( int initialValue = 0 ) {}
    int read() const { return _value; }
    void write( int newValue ) { _value = newValue; }

private:
    int _value;
};

#endif
```

1.4.2 IntCell implementation

```
// IntCell.cpp
#include "IntCell.h" // we're making use IntCell's interf. so we include it.

IntCell::IntCell( int initialValue = ) : _value { initialValue } { }

// Return the stored value.
int IntCell::read() const { return _value; }

// Change the stored value.
void IntCell::write( int newValue ) { _value = newValue; }

};
```

Notice that we must use the **scope resolution operator** (`::`) to identify the class each member function belongs to. Otherwise, it's assumed that the function is in the **global scope**. Also note that the signature of an implemented member function *must* match exactly the signature in the class interface.

1.4.3 Using IntCell

```
// main.cpp
#include "IntCell.h"

int main() {
    IntCell m{};
    m.write(5);
    int val = m.read(); // 5

    return 0;
}
```

To compile the program, run:

```
g++ -o prog main.cpp IntCell.cpp
```

The header files (e.g., `IntCell.h`) aren't listed since they're included in the implementation files and thus are compiled too.

1.5 C++ Objects

In C++ objects are declared much like primitive types (e.g., `char`, `int`, `float`, etc.):

```
IntCell obj1;      // Zero parameter constructor
IntCell obj2(12);  // One parameter constructor
```

However, due to an effort to standardize the uniform initialization syntax using braces, it's recommended to declare objects as follows:

```
IntCell obj1;      // Zero parameter constructor
IntCell obj1{};    // same as before
IntCell obj2{12};  // One parameter constructor
```

1.6 vectors and strings

The C++ standard defines the classes `vector` and `string` that intend to replace the built-in C++ array. Unlike arrays, these classes behave like first-class objects. Thus, they can be copied with `=`, they remember how many items they can store, and their indexing operator check that the provided valid (to access some element) is valid.

1.6.1 vector

```
vector<int> numbers = {1, 2, 3, 4, 5}; // Or...
vector<int> numbers {1, 2, 3, 4, 5};
```

However, there's some ambiguity with the declaration. For instance, `vector<int> a(12)` declares a vector that stores 12 integers but `vector<int> a{12}` declares a vector of size 1 with the value 12 at position 0. This is because C++ gives precedence to the initializer list. If the intention is to declare a vector of size 12, use the old C++ syntax using parentheses: `vector<int> a(12)`.

Instead of using the regular for loop, vectors can be looped over using the following syntax:

```
int sum = 0;
vector<int> numbers = {1, 2, 3, 4, 5};
for( int num : numbers ) {
    sum += num;
}
```

The keyword `auto` can be used to let the compiler determine the type:

```
int sum = 0;
vector<int> numbers = {1, 2, 3, 4, 5};
for( auto num : numbers ) {
    sum += num;
}
```

However, keep in mind that this range syntax is only appropriate when 1) accessing elements sequentially and 2) when the index isn't needed.

1.7 C++ details

1.7.1 Pointers

A **pointer** is a variable that stores the address of a memory location (i.e., where an object resides). The syntax for declaring a pointer is as follows:

```
char* ptr_c;    // ptr_c is a pointer-to-char
int *ptr_i;     // ptr_i is a pointer-to-int
float *ptr_f;   // ptr_f is a pointer-to-float
void *u_ptr;    // ptr_u is a untyped pointer
```

Note that the position of the asterisk `*` (known as the **indirection operator** in this context) doesn't matter.

In order to obtain an object's address, the **address-of** operator & is used and a pointer variable is used to store it:

```
int num = 5;
int *ptr_num = &num; // get the address of num
int val = *ptr_num;   // get the value stored in the address ptr_num points
                      // to
*ptr_num = 6;         // now num is also 6. after all ptr_num to the same
                      // memory location
```

Thus, the indirection operator has two jobs: 1) it declares a pointer and 2) it dereferences a pointer.

1.7.2 Dynamic object allocation and garbage collection

Objects are created dynamically by allocating memory with **new** which returns a pointer to the newly created object:

```
IntCell *m;
m = new IntCell(); // OK
m = new IntCell{}; // OK
m = new IntCell;   // OK
```

If a pointer variable points at a class type, then a (visible) member of the object being pointed at can be accessed via the **->** operator:

```
int x = m->read();
m->write(6);
```

When an object that is allocated by **new** is no longer referenced, the **delete** operation must be applied to the object (through a pointer). Otherwise, the memory that it consumes is lost (until the program terminates). This is known as a **memory leak**.

```
delete m;
```

1.7.3 Lvalues, and Rvalues

An **lvalue** is an expression that identifies a non-temporary object:

```
vector<int> arr(3);
const int x = 2;
int y;
int z = x + y;
vector<string> *ptr = &arr;
```

The expressions `arr`, `x`, `y`, `z`, `ptr`, `*ptr`, and `z` are all lvalues.

An **rvalue** is an expression that identifies 1) a temporary object or 2) a value (e.g., literal constant) not associated with any object:

```
vector<int> arr(3);
const int x = 2;
int y;
int z = x + y;
vector<string> *ptr = &arr;
```

Here, `2`, `x+y`, and `&arr` are all rvalues.

1.7.4 References

A reference type allows us to define a new name for an existing value. In classic C++, a reference can generally only be a name for an lvalue, since having a reference to a temporary would lead to the ability to access an object that has theoretically been declared as no longer needed, and thus may have had its resources reclaimed for another object. However, in C++11, we can two types of references:

- An **lvalue reference** is declared by placing an `&` after some type. An lvalue reference then becomes a synonym for the object it references.

```
string str = "hello";
string &str_r = str;
str_r += " world"; // now str is "hello world"

string &lit_r = "hi"; // ILLEGAL: "hi" is not modifiable
string &bad_r = str + "!"; // ILLEGAL: str + "!" isn't an lvalue
```

- An **rvalue reference** is declared by placing an `&&` after some type. An rvalue reference has the same characteristics as an lvalue reference except that, unlike an lvalue reference, *an rvalue reference can also reference an rvalue* (i.e., a temporary).

```
string str = "hell";
string && bad1 = "hello"; // LEGAL
string && bad2 = str + ""; // LEGAL
string && sub = str.substr( 0, 4 ); // LEGAL
```

1.7.5 Uses of lvalue references

- **Aliasing complicated names.** We can rename objects that are too long and complicated to simpler names.

```

auto &whichList = theLists[ myhash( x, theLists.size( ) ) ];
if( find( begin( whichList ), end( whichList ), x ) != end( whichList ) )
{
    return false;
}
whichList.push_back( x );

```

Note that simply writing `auto whichList = theLists[myhash(x, theLists.size())];` wouldn't work because this would create a copy, and `whichList.push_back(x);` would be applied to the copy, not the original.

- **Range for loops.** By default, a range **for** loop cannot change the elements it iterates over, however taking a lvalue reference allows to modify those elements.

```

vector<int> numbers = {1, 2, 3};

for( auto &number : numbers ) {
    number++;
}

```

- **Avoiding a copy.** Given a function that returns an element of an array/vector, we could return a non-modifiable reference to that element instead of a copy. For instance, instead of

```

// @arr: A non-empty vector of strings.
// @return the maximum string in the @arr.
// Will abort() if @arr is empty.
string FindMax1(constvector<string>&arr) {
    if (arr.empty()) abort();
    int max_index=0;
    for (inti =1; i < arr.size(); ++i) {
        if(arr[max_index]<arr[i]) { max_index = i; }
    }
    return arr[max_index];
}

```

we could have

```

// @arr: A non-empty vector of strings.
// @return the maximum string in the @arr.
// Will abort() if @arr is empty.
const string &FindMax1(constvector<string>&arr) {
    if (arr.empty()) abort();
    int max_index=0;
    for (inti =1; i < arr.size(); ++i) {
        if(arr[max_index]<arr[i]) { max_index = i; }
    }
    return arr[max_index];
}

```

Syntax is needed in function declarations and returns to enable the passing and returning using references instead of copies. Notice the **const** keyword and & in the function's header.

2 Lecture 2: C++ Basics (cont.)

2.1 Parameter passing

Many languages, including C++, pass all parameters using **call-by-value**: the actual argument is copied into the formal parameter. However, this might be utterly inefficient if large complex objects are being passed since they're copied in their entirety. Historically C++ has had three ways to pass parameters:

- **call-by-value**: Useful to pass small objects that shouldn't be mutated by the function.

```
double average( double a, double b );
double x = 2.5, y = 3.5;
double z = average(x, y); // x and y remain unchanged

void swap( double a, double b ); // swap a and b
swap(x, y); // Not what's expected; x and y remain unchanged
```

- **call-by-reference (call-by-lvalue-reference)**: Useful for all type of objects that may be changed by the function.

```
double x = 2.5, y = 3.5;
swap( &a, &b );
swap(x, y); // now x = 3.5 and y = 2.5
```

- **call-by-constant-reference**: Useful for large objects that are expensive to copy and that must not be changed by the function. For this, the parameter is declared as a reference and the keyword **const** is used to signal that it cannot be modified.

```
string randomItem( const vector<string> &arr ); // return a random item in
arr
```

There's still another fourth way to pass parameters:

- **call-by-rvalue-reference**: Instead of copying a temporary object stored in an rvalue, a move is used; moving an object's state is easier than copying, as *it may involve just a simpler pointer change*. Functions know if a value is temporary or not based on their signature so the primary use for this type of parameter passing is overloading a function based on whether a parameter is an lvalue or rvalue.

```

string randomItem( const vector<string> &arr ); // return random item in
    lvalue arr (single &)
string randomItem(      vector<string> &&arr ); // return random item in
    rvalue arr (double &)

vector<string> v = {"hello", "world"};
string x = randomItem(v);                // invoke lvalue function
string x = randomItem({"hello", "world"}); // invoke rvalue function

```

This idiom is particularly useful for defining the behavior of = and in writing constructors.

2.2 Return passing

In C++, there are several mechanism for returning from a function:

- **return-by-value:** A copy of the object is returned which can be inefficient, however in C++11 return-by-value may still be efficient for large objects if the returned objects are rvalues.

```

double average( double a, double b );
LargeType randomItem( const vector<LargeType> &arr ); // potentially
    inefficient
vector<int> partialSum( const vector<int> &arr );      // efficient in C
    ++11

```

The following are two versions of the function `randomItem`. The second version avoids the creation of a temporary `LargeType` object, but only if the caller accesses it with a constant reference:

```

LargeType randomItem1( const vector<LargeType> &arr ) {
    return arr[ randomInt(0, arr.size() -1) ];
}

const LargeType & randomItem2( const vector<LargeType> &arr ) {
    return arr[ randomInt(0, arr.size() -1) ];
}

vector<LargeType> vec;
LargeType item1 = randomItem2(vec);          // copy
LargeType item2 = randomItem2(vec);          // copy
const LargeType &item3 = randomItem2(vec);    // no copy
auto &item4 = randomItem2(vec);              // no copy

```

- **return-by-constant-reference:** This avoid creating an immediate copy of an object. However, the caller must also use a constant reference to access the return value; otherwise, there will be still a copy. What does the constant reference mean? It means that we don't want to allow changes to be made by the caller by using the return value.

```
const LargeType & randomItem2( const vector<LargeType> &arr ) {  
    return arr[ randomInt(0, arr.size() -1) ];  
}  
  
vector<LargeType> vec;  
const LargeType &item3 = randomItem2(vec); // no copy
```

- **return-by-reference:** A reference is returned and the caller can modify the returned value. This is used in a few places to allow the caller of a function to have modifiable access to the internal data representation of a class.

2.3 C++11 `std::swap` and `std::move`

Given that copying large objects is expensive, C++11 allows the programmer to easily replace expensive copies with moves provided the object's class supports moves.

Take the following example of a `swap` function that swap its arguments by three copies:

```
void swap( vector<string> &x, vector<string> &y ) {  
    vector<string> tmp = x;  
    x = y;  
    y = tmp;  
}
```

In C++11, if the right-hand side of the assignment operator (or constructor) is an rvalue, then if the object supports moving, we can automatically avoid copies. In the example above, we know that `vector` supports moving so instead of copy operations we could do move operations. These could be done either by casting the right-hand side of an assignment to an rvalue reference or by using `std::move`.

```
// Using type-casting  
void swap( vector<string> &x, vector<string> &y ) {  
    vector<string> tmp = static_cast<vector<string>> &&>( x );  
    x = static_cast<vector<string> &&>( y );  
    y = static_cast<vector<string> &&>( tmp );  
}  
  
// Using std::move, equivalent to casting but more succinct  
void swap( vector<string> &x, vector<string> &y ) {  
    vector<string> tmp = std::move(x);  
    x = std::move(y);  
    y = std::move(tmp);  
}
```

NOTE: `std::move` doesn't move anything; rather, it makes a value (either lvalue or rvalue) subject to

be moved.

It's worth noting that `std::swap` is already part of STL and works for any part:

```
vector<string> x;  
vector<string> y;  
std::swap(x, y); // x contains y's contents and y contains x's contents
```

2.4 The big five

In C++, classes come with five special functions already written for each class. These are the **destructor**, **copy constructor**, **move constructor**, **copy assignment operator**, and **move assignment operator**. In many cases, you accept the default behavior provided by the compiler can be relied on. Sometimes you cannot.

Let's assume the following interface for the `IntCell` class:

```
#ifndef INTCELL_H  
#define INTCELL_H  
  
class IntCell {  
public:  
    explicit IntCell( int initialValue = 0 ) { value_ = new int{  
        initialValue}; }  
    int read() const {}  
    void write( int x ) {}  
private:  
    int *value_;  
};  
  
#endif
```

2.4.1 Destructor

This function is called whenever an object goes out of scope or is subjected to a `delete` operation. Typically, the only responsibility of the destructor is to free up any resources that were acquired during the use of the object. This includes calling `delete` for any corresponding `new`s, closing any files that were opened, and so on.

```
IntCell::~IntCell() {  
    delete value_;  
}
```

2.4.2 Copy constructor and move constructor

These two constructors are required to construct a new object, initialized to the same state as another object of the same type.

- If the existing object is an **lvalue**, then it's a **copy constructor**.
- If the existing object is an **rvalue**, then it's a **move constructor**.

```
// Copy constructor, the parameter is an lvalue of the same type.
IntCell::IntCell( const IntCell &rhs ) {
    stored_ = new int{*rhs.value_};
}

// Move constructor, the parameter is an rvalue of the same type.
IntCell::IntCell( IntCell &&rhs ) : value_{rhs.value_} {
    rhs.value_ = nullptr;
}
```

When is either constructor called?

- During a declaration with initialization.

```
IntCell B = C;    // Copy constructor if C is lvalue; Move constructor if C
                  is rvalue
IntCell B { C }; // same as above
```

- An object passed using call-by-value (rarely done).
- An object returned by value.

2.4.3 Copy assignment and move assignment (operator =)

The assignment operator is called when = is applied to two objects that have both been previously constructed. Given the expression `lhs = rhs`, then the state of `rhs` (right hand side) is copied into `lhs` (left hand side).

- If `rhs` is an **lvalue**, this is done using the copy assignment operator.
- If `rhs` is an **rvalue**, this is done using the move assignment operator.


```
// Copy assignment operator
IntCell & IntCell::operator=( const IntCell &rhs ) {
    if (this &= &rhs) {
        *value_ = *rhs.value_;
    }
    return *this;
}

// Move assignment operator
IntCell & IntCell::operator=( IntCell &&rhs ) {
    std::swap(value_, rhs.value_);
    return *this;
}
```

In C++11, the copy assignment operator could be implemented more idiomatically with a copy-and-swap idiom:

```
// Copy assignment operator
IntCell & IntCell::operator=( const IntCell &rhs ) {
    IntCell copy = rhs; // calls the copy constructor
    std::swap(*this, copy);
    return *this;
}
```

2.4.4 Defaults

It's often the case that the defaults “big five” are perfectly acceptable, so nothing need to be done. If a class consists of data members that are exclusively primitive types and objects for which the defaults make sense, the class defaults will usually make sense. Thus a class whose data members are **int**, **double**, **string**, and even **vector<string>** can accept defaults.

When do defaults fail? The defaults fail in a class that contains a data member that is a pointer:

- The default destructor does nothing to data members that are pointers. It's imperative that we **delete** them ourselves.
- The copy constructor and copy assignment operator both copy the value of the pointer (i.e., a memory address) rather than the object being pointed at. This means we end up with two class instances that contain pointers that point to the same objects being pointed at. This is known as **shallow copy**. However, we would typically expect a **deep copy**, in which a clone of the entire object is created.

Thus, as a result, when a class contains pointers as data members, and deep semantics are important, we typically must implement the destructor, copy assignment, and copy constructors ourselves. Doing so removes the move defaults, so we also must implement move assignment and the move constructor. As a general rule, either you accept the default for all five operations, or you should declare all five, and explicitly define, default (use the keyword **default**), or disallow each (use the keyword **delete**). Generally we will defined all five.

NOTE: If you write any of the big-five, it would be good practice to explicitly consider all the others, as the defaults may be invalid or inappropriate. Changing a default method is an all or nothing situation.

When don't the defaults work? The most common situation in which the defaults do not work occurs when a data member is a pointer type and the pointer is allocated by some object member function (such as the constructor). A problem that might arise is the default copy assignment and copy constructor doing shallow copies. Another problem is **memory leak**; an object allocated by a constructor might remain allocated and it's never reclaimed.

2.5 Templates

An algorithm is **type independent** if the logic of the algorithm does not depend on the type of items it handles. When we write C++ code for a type-independent algorithm or data structure, it's preferable to write the code once rather than recode it for each different type. This is accomplished by using **templates**.

2.5.1 Function templates

A **function template** isn't an actual function, but instead a pattern for what could become a function. For example, the following is a function template:

```
/*
Return the maximum item in array a.
Assumes a.size() > 0.
Comparable object must provide operator< and operator=.
*/

template <typename Comparable>
const Comparable & findMax( const vector<Comparable> &a ) {
    int maxIndex = 0;
    for (int i = 1; i < a.size(); i++) {
        if (a[maxIndex] < a[i]) {
            maxIndex = i;
        }
    }
    return a[maxIndex];
}
```

Using the function template would look as follows:

```
vector<int> v1 = { 37, 12, 1, 89 };
vector<double> v2 = { 78.1, 89.8, 12.4, 1.1 };
vector<string> v3 = { 'hi', 'ha', 'ho', 'he' };
vector<IntCell> v4(75);

findMax(v1); // OK: Comparable = int
findMax(v2); // OK: Comparable = double
findMax(v3); // OK: Comparable = string
findMax(v4); // Illegal; IntCell doesn't implement operator< and thus not
               a Comparable
```

Thus, one of the limitations of function templates is that object passed to the function need to implement whatever operator the function uses internally. However, instead of relying on “hardcoded” conditions, a function that perform the required operation could be passed alongside the object. This type of functions are known as **function objects**.

2.6 Function objects

Previously we implemented `findMax` as a function template but it was only limited to objects that have an `operator<` function defined. Instead, we need to rewrite `findMax` to accept as parameters an array of object and a comparison function that explains how to decide which of two objects is the larger and which is the smaller. Instead of relying on the array objects knowing how to compare themselves, we completely decouple this information from the object in the arrays.

How do we pass a functions as parameters though? One way is to define a class with no data and one member function, and pass an instance of the class. In fact, the function is passed by placing it

inside an object, which is known as a **function object**.

The following is the simplest implementation of the function object idea:

```
#include <iostream>
#include <vector>
#include <string>

// Generic findMax, with a function object.
// Precondition: a.size() > 0.
// Comparator object is assumed to implement the isLessThan method.
template <typename Object, typename Comparator>
const Object & findMax( const std::vector<Object> &arr, Comparator cmp ) {
    int maxIndex = 0;

    for (int i = 1; i < arr.size(); i++) {
        if (cmp.isLessThan(arr[maxIndex], arr[i])) {
            maxIndex = i;
        }
    }

    return arr[maxIndex];
};

class StringComparisonByLength {
public:
    bool isLessThan( const std::string &lhs, const std::string &rhs )
    {
        return lhs.length() < rhs.length();
    }
};

class StringComparisonCaseInsensitive {
public:
    bool isLessThan( const std::string &lhs, const std::string &rhs )
    {
        return std::strcasecmp(lhs.c_str(), rhs.c_str()) < 0;
    }
};

int main() {
    std::vector<std::string> greetings = {"hi", "hello", "bonjour", "HOLA"};

    std::cout << findMax( greetings, StringComparisonByLength{} ) << "\n";
    std::cout << findMax( greetings, StringComparisonCaseInsensitive{} )
        << "\n";

    return 0;
}
```

C++ function objects are implemented using this basic idea, but with some fancy syntax. First, instead of using a function with a name, we use operator overloading. Instead of the function being `isLessThan`, it is `operator()`. Second, when invoking `operator()`, `cmp.operator()(x,y)` can be shortened to `cmp(x,y)`. Third, we can provide a version of `findMax` that works without a function object that uses a default ordering; the implementation uses the Standard Library function object template `less` (defined in header file `functional`) to generate a function object that imposes the normal default ordering.

The following is the more idiomatic implementation:

```
#include <iostream>
#include <vector>
#include <string>

// Generic findMax, with a function object.
// Precondition: a.size() > 0.
template <typename Object, typename Comparator>
const Object & findMax( const std::vector<Object> &arr, Comparator
    isLessThan ) {
    int maxIndex = 0;

    for (int i = 1; i < arr.size(); i++) {
        if (isLessThan(arr[maxIndex], arr[i])) {
            maxIndex = i;
        }
    }

    return arr[maxIndex];
};

// Generic findMax, using default ordering.
const Object & findMax( const vector<Object> &arr ) {
    return findMax(arr, less<Object>{} );
}

class StringComparisonByLength {
public:
    bool operator()( const std::string &lhs, const std::string &rhs )
    {
        return lhs.length() < rhs.length();
    }
};

class StringComparisonCaseInsensitive {
public:
    bool operator()( const std::string &lhs, const std::string &rhs )
    {
        return std::strcasecmp(lhs.c_str(), rhs.c_str()) < 0;
    }
};

int main() {
    std::vector<std::string> greetings = {"hi", "hello", "bonjour", "HOLA"};

    std::cout << findMax( greetings, StringComparisonByLength{} ) << "\n";
    std::cout << findMax( greetings, StringComparisonCaseInsensitive{} )
        << "\n";
    std::cout << findMax( greetings ) << "\n";

    return 0;
}
```

2.6.1 Class templates

In its most simplest form, a class template works much like a function template. The following `MultiCell` is an implementation that is like `IntCell`, but works for any type `Object`, provided that `Object` has a zero-parameter constructor, a copy constructor, and a copy assignment operator.

```
template <typename Object>
class MemoryCell {
public:
    explicit MemoryCell( const Object &initialValue = Object{} ) :
        _value{initialValue} {}
    const Object & read() const { return _value; }
    void write( const Object &x ) { _value = x; }

private:
    Object _value;
};
```

Notice that

- `Object` is passed by constant reference.
- the default parameter for the constructor is not 0, because 0 might not be a valid `Object`. Instead, the default parameter is the result of constructing an `Object` with its zero-parameter constructor.

If we implement class templates as a single unit, then there is very little syntax baggage. Many class templates are, in fact, implemented this way because, currently, separate compilation of templates does not work well on many platforms. **Therefore, in many cases, the entire class, with its implementation, must be placed in a .h file.** Popular implementations of the STL follow this strategy.

2.7 Using matrices

This will be implemented by using a vector of vectors (e.g., a vector of `int` vectors).

```
#ifndef MATRIX_H
#define MATRIX_H

#include <vector>

template <typename Object>
class Matrix {
public:
    /*
     * Create _array as having rows entries each of type vector<Object>
     * .
     * We have a rows zero-length vectors of Object so each row is
     * resized
     * to have cols columns. Thus this creates a two-dimensional array
     * .
     */
    Matrix( int rows, int cols ) : _array(rows) {
        for (auto &row : _array) {
            row.resize(cols);
        }
    }

    Matrix( std::vector<std::vector<Object>> v ) : _array{ v } { }

    Matrix( std::vector<std::vector<Object>> &&v ) : _array{ std::move
        (v) } { }

    /*
     * Array indexing operator. This returns the row (a vector<Object>
     * at
     * the index row.
     */
    const std::vector<Object> & operator[]( int row ) const {
        return _array[row]
    }

    int numrows() const {
        return _array.size();
    }

    int numcols() const {
        return numrows() ? _array[0].size() : 0;
    }

private:
    /*
     * A matrix is represented by _array, a vector of vector<Object>.
     */
    std::vector<std::vector<Object>> _array;
};

#endif
```


2.8 Class problems

1. What's $1 + 2 + 3 + \dots + n$? Prove it.

$$1 + 2 + 3 + \dots + n = \frac{n(n+1)}{2}$$

Proof by induction:

- **Base case:** When $n = 1$, then $1 = 1(1 + 1)/2 = 1$.
- **Inductive step:** Since $1 = (1 * 2)/2$, the statement $P(1)$ is true. Assume that $P(k)$ is true for an arbitrary positive integer k . We show that $P(k + 1)$ is true. In other words,

$$1 + 2 + 3 + \dots + (k + 1) = \frac{(k + 1)(k + 2)}{2}$$

Thus

$$\begin{aligned} 1 + 2 + 3 + \dots + (k + 1) &= (1 + 2 + 3 + \dots + k) + (k + 1) \\ &= \frac{k(k + 1)}{2} + (k + 1) \\ &= \frac{k(k + 1)}{2} + \frac{2(k + 1)}{2} \\ &= \frac{(k + 1)(k + 2)}{2} \end{aligned}$$

Therefore, by the principle of mathematical induction, $P(n)$ is true for every positive integer n .

2. What's $1 + 2 + 4 + \dots + 2^n$? Prove it.

$$\begin{aligned} n = 0 &\Rightarrow 1 \Rightarrow 1 \\ n = 1 &\Rightarrow 1 + 2 \Rightarrow 3 \\ n = 2 &\Rightarrow 1 + 2 + 4 \Rightarrow 7 \\ n = 3 &\Rightarrow 1 + 2 + 4 + 8 \Rightarrow 15 \\ &\dots \\ n = n &\Rightarrow 1 + 2 + \dots + 2^n \Rightarrow 2^{(n+1)} - 1 \end{aligned}$$

Thus, $1 + 2 + 4 + \dots + 2^n = 2^{(n+1)} - 1$.

Proof by induction:

- **Base case:** When $n = 0$, then $1 = 2^{(0+1)} - 1 = 1$. When $n = 1$, then $2^0 + 2^1 = 3 = 2^{(1+1)} - 1 = 3$.
- **Inductive step:** Since $1 = 2^{(0+1)} - 1$, the statement $P(1)$ is true. Assume that $P(k)$ is true for an arbitrary positive integer k . We show that $P(k + 1)$ is true. In other words, we show that

$$1 + 2 + 4 + \dots + 2^{(k+1)} = 2^{(k+2)} - 1.$$

Thus,

$$\begin{aligned} 1 + 2 + 4 + \dots + 2^{(k+1)} &= (1 + 2 + 4 + \dots + 2^k) + 2^{(k+1)} \\ &= 2^{(k+1)} - 1 + 2^{(k+1)} \\ &= 2^{(k+1)} + 2^{(k+1)} - 1 \\ &= 2 \cdot 2^{(k+1)} - 1 \\ &= 2^{(k+2)} - 1 \end{aligned}$$

Therefore, by the principle of mathematical induction, $P(n)$ is true for every non-negative integer n .

3. If $A_0 = 1$ and $A_n = 2A_{(n-1)} + 1$, what's a closed formula for A_n ? Prove it.

$$A_1 = 2 \cdot A_0 + 1 = 3$$

$$A_2 = 2 \cdot A_1 + 1 = 7$$

$$A_3 = 2 \cdot A_2 + 1 = 15$$

$$A_4 = 2 \cdot A_3 + 1 = 31$$

...

$$A_n = 2 \cdot A_{(n-1)} + 1 = 2^{(n+1)} - 1$$

3 Lecture 3: Algorithm analysis

An **algorithm** is a clearly defined set of simple instructions which must be followed in order to solve a particular problem.

3.1 Mathematical background

Algorithm analysis is grounded on mathematics and the definitions below set up a formal framework to study algorithms. These definitions try to establish a **relative order among functions**. Given two functions, there are usually points where one function is smaller than the other so it doesn't make sense to claim, for instance, $f(N) < g(N)$. Instead, we compare their **relative rates of growth**.

Big-Oh $T(N) = O(f(N))$ if there are positive constants c and n_0 such that $T(N) \leq c \cdot f(N)$ when $N \geq n_0$.

Informally, the growth rate $T(N)$ is less than or equal to that $f(N)$.

Big-Omega $T(N) = \Omega(g(N))$ if there are positive constants c and n_0 such that $T(N) \geq c \cdot g(N)$ when $N \geq n_0$.

Informally, the growth rate $T(N)$ is greater than or equal to that $g(N)$.

Big-Theta $T(N) = \Theta(h(N))$ **if and only if** $T(N) = O(h(N))$ and $T(N) = \Omega(h(N))$.

Informally, the growth rate $T(N)$ equals the growth rate of $h(N)$. This means that $T(N)$ is eventually *sandwiched* between two different constant multiples of $h(N)$.

Little-Oh $T(N) = o(p(N))$ if, for all positive constants c , there exist an n_0 such that $T(N) < c \cdot p(N)$ when $N > n_0$.

Informally, the growth rate $T(N)$ is less than the growth rate of $f(N)$. This is different from Big-Oh, given that Big-Oh allows the possibility that the growth rates are the same.

NOTE: In the above definitions, both c and n_0 are constants and they cannot depend on N . If you see yourself saying "take $n_0 = N$ " or "take $c = \log_2 N$ " in an alleged Big-Oh proof, then you need to start with choices of c and n_0 that are independent of N .

3.2 Notation

3.2.1 Big-Oh notation

When we say that $T(N) = O(f(N))$, we're guaranteeing that the function $T(N)$ grows at a rate no faster than $f(N)$; thus $f(N)$ is an **upper bound** on $T(N)$. Alternatively, we could say that $T(N)$ is

bounded above by a constant multiple of $f(N)$.





Figure 2.1: A picture illustrating when $T(n) = O(f(n))$. The constant c quantifies the “constant multiple” of $f(n)$, and the constant n_0 quantifies “eventually.”

Figure 1: Big-Oh

For example, $T(N) = O(N^2)$ means that the function $T(N)$ grows at a rate no faster than N^2 ; its growth could equal that of N^2 but it could never surpass it.

3.2.2 Big-Omega notation

When we say that $T(N) = \Omega(g(N))$, we’re guaranteeing that the function $T(N)$ grows at a rate no lower than $g(N)$; thus $g(N)$ is a **lower bound** on $T(N)$. Alternatively, we could say that $T(N)$ is **bounded below** by constant mutiple of $g(N)$.



$T(n)$ again corresponds to the function with the solid line. The function $f(n)$ is the upper dashed line. This function does not bound $T(n)$ from below, but if we multiply it by the constant $c = \frac{1}{4}$, the result (the lower dashed line) does bound $T(n)$ from below for all n past the crossover point at n_0 . Thus $T(n) = \Omega(f(n))$.

Figure 2: Big-Omega

For example, $T(N) = \Omega(N^2)$ means that the function $T(N)$ grows at a rate no lower than N^2 ; its growth could equal that of N^2 but it could never drop below it.

Whenever talking about Big-Oh, there's always an implication about Big-Omega. For instance, $T(N) = O(f(N))$ (i.e., $f(N)$ is an **upper bound** on $T(N)$) implies that $f(N) = \Omega(T(N))$ (i.e., $T(N)$ is a **lower bound** on $f(N)$). As an example, N^3 grows faster than N^2 , so we can say that $N^2 = O(N^3)$ or $N^3 = \Omega(N^2)$.

3.2.3 Big-Theta notation

When we say that $T(N) = \Theta(h(N))$, we're guaranteeing that the function $T(N)$ grows at the same rate as $h(N)$. However, when two functions grow at the same rate, then the decision of whether or not to signify this with $\Theta()$ can depend on the context.

3.3 Typical growth rates

Function	Name
c	Constant
$\log_2 N$	Logarithmic
$\log^2 N$	Log-squared
N	Linear
$N \log N$	
N^2	Quadratic
N^3	Cubic
2^N	Exponential

3.4 Rules

Rule 1 If $T_1(N) = O(f(N))$ and $T_2(N) = O(g(N))$, then

1. $T_1(N) + T_2(N) = O(f(N) + g(N))$. Less formally it is $O(\max(f(N), g(N)))$, and
2. $T_1(N) \cdot T_2(N) = O(f(N) \cdot g(N))$

Rule 2 If $T(N)$ is a polynomial of degree k , then $T(N) = \Theta(N^k)$.

Rule 3 $\log^k N = O(N)$ for any constant k . In other words, logarithms run at a lower rate than linear functions which means logarithms grow very slowly.

3.5 Few points

- In simple terms, the goal of asymptotic notation is *to suppress constants factors and lower-order*. Thus, they aren't included in Big-Oh. For instance, don't say $T(N) = O(2N^2)$ or $T(N) = O(N^2 + N)$. In both cases, the correct form is $T(N) = O(N^2)$.

When analyzing the running time of an algorithm, why would we want to throw away information like constant factors and lower-order terms? 1) Lower-order terms become increasingly irrelevant as you focus in large inputs, which are the inputs that require algorithmic ingenuity and 2) constant factors are generally highly dependent on the details of the environment (e.g., programming language, architecture, compiler, etc.) and thus ditching them allows to generalize by not committing ourselves to a specific programming language, architecture, etc.

- The relative growth rates of two functions $f(N)$ and $g(N)$ can always be determined by computing $\lim_{N \rightarrow \infty} \frac{f(N)}{g(N)}$, using [L'Hopital's rule](#) if necessary. The limit can have four possible values:

- The limit is 0, meaning that $f(N) = O(g(N))$.
 - The limit is $c \neq 0$, meaning that $f(N) = \Theta(g(N))$.
 - The limit is ∞ , meaning that $g(N) = o(f(N))$.
 - The limit doesn't exist and thus no relation exist.
- Stylistically it's bad to say that $f(N) \leq O(g(N))$ since the inequality is already implied by the definition of Big-Oh.
 - It's wrong to write $f(N) \geq O(g(N))$, because it doesn't make sense.

3.6 Model of computation

The model of computation is basically a normal computer with the following characteristics:

- Instructions are executed sequentially.
- The model has the standard repertoires of simple instructions, such as addition, multiplication, comparison, and assignment. Unlike real computers, this computer takes exactly one time unit to do anything.
- The computer has fixed-size (e.g., 32-bit) integers and no fancy operations (e.g., matrix inversion, sorting, etc.).
- The computer has infinite memory.

3.7 What to analyze

The most important resource to analyze is generally the *running time* and typically, the size of the input is the main consideration. We define two functions, $T_{avg}(N)$ (for the average-case running time) and $T_{worst}(N)$ (for the worst-case running time) used by an algorithm on input of size N . It's worth noting that $T_{avg}(N) \leq T_{worst}(N)$ (i.e., $T_{avg}(N)$ has a lower rate growth than that of $T_{worst}(N)$).

3.7.1 Types of performance

Best-case performance Although it can be occasionally analyzed, it's often of little interest since it doesn't represent typical behavior.

Average-case performance It often reflects typical behavior, however it doesn't provide a bound for all input and it can be difficult to compute. Furthermore, it's also hard to define what's the average input

Worst-case performance It represents a guarantee for performance on any possible input. This is the quantity generally required because it provides a bound for all input.

3.8 Running-time calculations

When computing a Big-Oh running time, there are several general rules:

For loops The running time of a **for** loop is *at most* the running time of the statements inside the **for** loop times the number of iterations.

Nested loops Analyze these inside out. The total running time of a statement inside a group of nested loops is the running time of the statement multiplied by the product of the sizes of all the loops.

For example, the following program is $O(N^2)$ because the inner loop does N iterations and the outer loops does N too. This amounts to $N \times N = N^2$ iterations. The first assignment counts as one operation (thus $O(N)$) while the innermost statement counts for 2 operations (1 multiplication and 1 assignment). Thus, to be more precise the program's running time is $1 + 2N^2$, however Big-Oh suppresses constant factors and lower-order terms.

```
k ← 0
for [1, n] → i:
  for [1, n] → j:
    k ← i * j
```

Consecutive statements They just add (which means that the maximum is the one that counts).

For example, the following program fragment, which has $O(N)$ work followed by $O(N^2)$ work, is ultimately $O(N^2)$ which dominates $O(N)$:

```
for [0, n) → i:
  a[i] = 0

for [0, n) → i:
  for [0, n) → j:
    a[i] += a[j] + i + j
```

If/else For the fragment **if** condition { S1 } **else** { S2 }, the running time of an **if/else** statement is never more than the running time of the test plus the larger of the running times of S1 and S2.

3.9 Sample problems

3.9.1 Sum of cubes

Input: a positive integer n .

Output: the sum of all cubes from 1 to n^3 .

```
SumOfCubes(n):  
    sum ← 0                # 0(1)  
    for [1, n] → i:        # 0(n)  
        sum += i * i * i    # 0(4), 1 assignment, 1 addition and 2 products  
    return sum              # 0(1)
```

Thus, for $\text{SumOfCubes}(N) = 1 + 4n + 1 = 4n + 2 = O(n)$. We discard both the constant factors and the lower-order terms.

3.9.2 Factorial

The factorial is defined as

$$n! = n \times (n-1) \times (n-2) \times \cdots \times 2 \times 1$$

Implemented recursively, the algorithm is as follows:

Input: a non-negative integer n

Output: the factorial of n

```
Factorial(n):  
    if n ≤ 1:  
        return 1  
    else:  
        return n * Factorial(n-1)
```

However, this is a thinly veiled **for** loop. In this case, the analysis involves the recurrence relation $T(n) = 1 + T(n-1)$ for $n > 1$, $T(1) = 2$ that needs to be solved:

$$\begin{aligned}
T(N) &= 1 + T(n-1) \\
&= 1 + (1 + T(n-2)) \\
&= 1 + (1 + (1 + T(n-3))) \\
&= \dots \\
&= 1 + (1 + (1 + T(n-k))) \text{ } k \text{ 1's}
\end{aligned}$$

$$\begin{aligned}
T(N) &= k + T(n-k) \\
&= (n-1) + T(n-(n-1)) \\
&= (n-1) + T(1) = (n-1) + 2 = n + 1
\end{aligned}$$

Thus, $T(n) = O(n)$.

3.9.3 Maximum subsequence sum problem

Given (possibly negative) integers A_1, A_2, \dots, A_N , find the maximum value of $\sum_{k=1}^i A_k$.

In other words, given a one-dimensional array of numbers, find a contiguous subarray with the largest sum. For example, with the input $-2, 11, -4, 13, -5, -2$ the answer is $11 - 4 + 13 = 20$.

This problem is interesting mainly because there are many algorithms to solve it, and the performance of these algorithms varies drastically.




Figure 3: Plot (N vs. time) of various algorithms

For algorithm 4 (linear), as the problem size increases by a factor of 10, so does the running time. For algorithm 3 (quadratic), a tenfold increase in input size yields roughly hundredfold (10^2) increase in running time.

3.9.4 Algorithm 1 (brute force)

implementations/max-subsequence-sum:

Input: array A of integers.
Output: the maximum positive subsequence sum.

```

MaxSubsequenceSum(A):
    maxSum ← 0
    for [0, A.size) → i:
        for [i, A.size) → j:
            currentSum ← 0
            for [i, j] → k:
                currentSum += A[k]
            if currentSum > maxSum:
                maxSum = currentSum
    return maxSum

```

For this algorithm, there's N starting places, average $\frac{N}{2}$ lengths to check, and average $\frac{N}{4}$ numbers to add. Thus we have $O(N^3)$.

3.9.5 Algorithm 2 (brute force)

[implementations/max-subsequence-sum:](#)

```
Input: array A of integers.
Output: the maximum positive subsequence sum.

MaxSubsequenceSum( A ):
    maxSum ← 0
    for [0, A.size) → i:
        currentSum ← 0
        for [i, A.size) → j:
            currentSum += A[j]
            if currentSum > maxSum:
                maxSum = currentSum
    return maxSum
```

For this algorithm, the innermost **for** loop has been removed. There's N starting places and an average $\frac{N}{4}$ numbers to add. Thus we have $O(N^2)$.

3.9.6 Algorithm 3

[implementations/max-subsequence-sum:](#)

Input: array A of integers.

Output: find the maximum sum **in** subarray spanning A[LEFT..RIGHT]. It does not attempt to maintain actual best sequence.

```

MaxSumRec( A, LEFT, RIGHT ):
    # Base case
    if LEFT = RIGHT:
        if A[LEFT] > 0:
            return A[LEFT]
        else
            return 0

    center ← (LEFT + RIGHT) div 2
    maxLeftSum ← MaxSumRec(A, LEFT, center)
    maxRightSum ← MaxSumRec(A, center + 1, RIGHT)

    maxLeftBorderSum ← 0
    leftBorderSum ← 0

    for [center, LEFT] → i:
        leftBorderSum += A[i]
        if leftBorderSum > maxLeftBorderSum:
            maxLeftBorderSum ← leftBorderSum

    maxRightBorderSum ← 0
    rightBorderSum ← 0

    for [center + 1, RIGHT] → j:
        rightBorderSum += A[j]
        if rightBorderSum > maxRightBorderSum:
            maxRightBorderSum ← rightBorderSum

    return max(maxLeftSum, maxRightSum, maxLeftBorderSum +
              maxRightBorderSum)

# Driver for divide-and-conquer maximum contiguous subsequence sum
algorithm.
MaxSubsequenceSum( A ):
    return MaxSumRec(A, 0, A.size - 1)

```

3.9.7 Algorithm 4

implementations/max-subsequence-sum:

Input: array A of integers.
Output: the maximum positive subsequence sum.

```
MaxSubsequenceSum( A ):  
    maxSum ← 0  
    currentSum ← 0  
  
    for [0, A.size] → j:  
        currentSum += A[j]  
  
        if currentSum > maxSum:  
            maxSum ← currentSum  
        else if currentSum < 0:  
            currentSum = 0  
  
    return maxSum
```

4 Lecture 4: Algorithm analysis (cont.) and lists/stacks/queues

4.1 Binary search

Given an integer X and integers A_0, A_1, \dots, A_{N-1} , which are **presorted** already in memory, find index i such that $A_i = X$, or return -1 if X is not in the group of integers.

implementations/binary-search

Input: array A of elements and item x which we are searching for.
Output: item where item is found or -1 if not found.

```
BinarySearch( A, x ):  
    low ← 0  
    high ← A.size - 1  
    while low ≤ high:  
        mid ← (low + high) / 2  
        if A[mid] < x:  
            low = mid + 1  
        else if A[mid] > x:  
            high = mid - 1  
        else:  
            return mid  
    return -1
```

4.1.1 Running time

We have the following recurrence relation $T(N) = 1 + T(N/2)$ and $T(1) = 1$. We must solve it to find out the algorithm's running time:

$$\begin{aligned}T(N) &= 1 + T(N/2) \\&= 1 + 1 + T(N/2^2) \\&= 1 + 1 + 1 + T(N/2^3) \\&\dots \\&= k + T(N/2^k)\end{aligned}$$

If N is a power of 2 (i.e., $N = 2^k$ with $k = \log(N)$) we'll have:

$$\begin{aligned}T(N) &= k + T(N/2^k) \\&= k + T(1) = \log(N) + 1\end{aligned}$$

This means that $T(N) = O(\log N)$.

4.2 List/stacks/queue ADTs

An **abstract data type** (ADT) is a set of objects (lists, sets, graphs, etc.) together with a set of related operations (add, remove, size, etc.). They provide a template for the objects, not how the objects and their respective set of operations are implemented.

4.2.1 The List ADT

We usually deal with a general list of the form $A_0, A_1, A_2, \dots, A_{N-1}$. We say that the size of this list is N . We will call the special list of size 0 an **empty list**. The **position** of the element A_i in a list is i .

The List ADT could describe the following operations:

find(x) Return the position of the occurrence of item x .

insert(x, i) Insert some element x at position i .

remove(i) Remove some element at position i .

4.2.2 Vector in the STL

Pros:

- Constant time indexing
- Fast to add data at the end (not the front).

Cons:

- Slow to add data in the middle.
- Inefficient for searches.

Operations

- `void push_back(const T &x)` - add `x` at the end of the list.
- `void pop_back()` - remove element at the end of the list.
- `const T &back() const` - return the element at the end of the list.
- `const T &front() const` - return the element at the front of the list.
- `void push_front(const T &x)` - add `x` to the front of the list.
- `void pop_front()` - remove the element at the front of the list.

4.2.3 List in the STL

Pros:

- Implemented as a doubly linked list.
- Fast insertion/removal of items in any position.

Cons:

- No indexing.
- Inefficient for searches.

Operations

- `T & operator[] (int idx)` - return element at index `idx` with no bounds checking.
- `T &at(int idx)` - return element at index `idx` with bounds checking.
- `int capacity() const` - return internal capacity of vector.
- `void reserve(int new_capacity)` - set new capacity and possibly void expansion of vector.

4.3 Containers

A **container** is a holder object that stores a collection of other objects. This is usually implemented as a class templates which provides it with great flexibility for the types supported as elements.

A container

- manages the storage space for its elements and provides member functions to access them either directly or through **iterators**.
- replicates structures commonly used in programming such as dynamic arrays (e.g., vectors), queues, stacks, heaps (e.g., priority queues), linked lists, trees (e.g., sets), associative arrays (e.g., maps), etc.

4.4 Iterators

Some operations on lists require the notion of a position. In the STL, a position is represented by some nested type known as an **iterator**. In particular, for a `list<string>`, the position is represented by the type `list<string>::iterator`; for a `vector<int>`, `Vector<int>::iterator`.

4.4.1 How to get an iterator

The STL lists (and all other STL containers) define a pair of methods:

- `iterator begin()`: returns an appropriate iterator representing the first item in the container.
- `iterator end()`: returns an appropriate iterator representing the endmarker in the container. This endmarker is “out-of-bounds”.

4.4.2 Iterator methods

- `itr++` and `++itr`: advances the iterator `itr` to the next location.
- `*itr`: returns a reference, which may or may not be modifiable, to the object stored at iterator `itr`'s location.
- `itr1 == itr2`: returns true if iterators `itr1` and `itr2` refer to the same location and false otherwise.
- `itr1 != itr2`: returns true if iterators `itr1` and `itr2` refer to a different location and false otherwise.

For example the code:

```
for (int i = 0; i <= v.size(); i++) {  
    std::cout << v[i] << "\n";  
}
```

could be rewritten as follows using iterators:

```
for (vector<int>::iterator itr = v.begin(); itr != v.end(); itr++) {  
    std::cout << *itr << "\n";  
}
```

Alternatively:

```
vector<int>::iterator itr = v.begin();  
while (itr != v.end()) {  
    std::cout << *itr++ << "\n";  
}
```

4.4.3 Container operations that require iterators

- `iterator insert(iterator pos, const T & x)`: adds `x` into the list, prior to the position by the iterator `pos`. This is a constant operator for `list`, but not for `vector`. The return value is an iterator representing the position of the inserted item.
- `iterator erase(iterator pos)`: removes the object at position and return the position of the element that followed `pos` prior to the call. It invalidates `pos`, making it stale.
- `iterator erase(iterator start, iterator pos)`: removes all items beginning at position up to, but not including `end`. An entire list `c` can be erased by `c.erase(c.begin(), c.end())`.

4.4.4 Example: Using erase on a list

```

/*
 * @brief Deletes every other element from lst, starting from the first
 *        item.
 * @param lst a list, or any object that supports iterators and erase.
 */
template <typename Container>
void removeEveryOtherItem( Container &lst ) {
    typename Container::iterator itr = lst.begin(); // could be shortened
    to                                              // auto itr = lst.
                                                    begin();

    while (itr != lst.end()) {
        itr = lst.erase(itr);
        if (itr != lst.end()) itr++;
    }
}

```

4.4.5 const_iterators

The result of `*itr` is both the value of the item that the iterator is viewing but also the item itself. This is very powerful but also introduces some complications.

Let's analyze the following routine that works for both `vector` and `list` and runs in linear time:

```

template <typename Container, typename Object>
void change( Container &c, const Object &newValue ) {
    typename Container::iterator itr = c.begin();
    while (itr != c.end()) {
        *itr = newValue; // set current iterator's value to newValue
        itr++;           // advance the pointer to next item
    }
}

```

The potential problem arises if `Container c` was passed to a routine using call-by-constant reference, meaning we would expect that not changes would be allowed to `c`, and the compiler would ensure this by not allowing calls to any `c`'s mutators. For example, consider the following code that prints a `list` of integers but also tries to do some changes:

```

void print( const list<int> &lst, ostream &out = cout ) {
    typename Container::iterator itr = lst.begin();
    while (itr != lst.end()) {
        out << *itr << "\n";
        *itr = 0; // <= This is the suspect.
        itr++;
    }
}

```

The solution provided by the STL is that every collection contains not only an `iterator` nested type but also `const_iterator` nested type. The main difference between them is that `operator*` for `const_iterator` returns a constant reference, and thus `itr*` for a `const_iterator` cannot appear on the left-hand side of an assignment statement.

Further the compiler will force you to use a `const_iterator` to traverse a *constant* collection and does so by providing two versions of `begin` and two versions of `end`:

- `iterator begin()`
- `const_iterator begin() const`
- `iterator end()`
- `const_iterator end() const`

Note: The two versions of `begin` can be in the same class only because of the const-ness of a method (whether an accessor or mutator) is considered to be part of the signature. The trick is overloading `operator[]`.

Using `auto` to declare your iterators means the compiler will deduce for you whether an `iterator` or `const_iterator` is substituted. This exempt the programmer from having to keep track of the correct iterator type and is precisely one of the intended uses of `auto`.

4.4.6 Printing a container

```
/*
 * @brief Prints out the container on the output stream.
 * @param c a given container.
 * @param out an output stream.
 */
template <typename Container>
void print( const Container &c, ostream &out = cout ) {
    if (c.empty()) {
        out << "(empty)";
    }
    else {
        typename Container::iterator itr = begin(c);
        out << "[" << *itr++; // print first item
        while (itr != end(c)) {
            out << ", " << *itr++;
        }
        out << "]" << "\n";
    }
}
```